

COMP 9500 research status

Vibhor Malik | Supervisor review | Updated 4 October 2026, Vancouver time

Experiment 1: a concrete, limited conclusion. Deterministic category enforcement works: TCP rises from 11.6% to 79.2%, and TCER falls from 6.1% to zero, with hard exclusion enforced by construction. Llama parses 120/120 templated requests but only 79/100 broader synthetic requests and 2/10 adversarial cases. The broader-set score is above the 70% floor and below the 85% target. All 10/10 recorded explanations equal the template. Boosting reduces NDCG@10 by 45.7%, and the non-inferiority check is uninformative at this low baseline. Backbone HR@1 is 0.0 to 0.46% across seeds, versus recency popularity's 8.0%. The study establishes bounded mechanical control; it does not establish an LLM advantage, greater naturalness or improved recommendation accuracy.

Experiment 1 decisions. The exposure hypothesis meets the reported directional tests, but identical rule-parser results prevent attributing the effect to the LLM. The 85% parser target fails. The absolute non-inferiority margin passes mechanically but permits complete baseline loss; the relative reading is unsupported. The recency comparison uses different information access and is descriptive. Experiment 1 remains frozen.

Experiment 2: pending, not a negative result. H1, higher intent satisfaction than both one-item feedback and the rule parser, and H2, greater intended movement per interaction, are not assessed. H3 quality differences are pending descriptive estimates with confidence intervals, without an NI claim. There are zero reserved ranking rows and no student agreement statistic. The bank contains 48 development and 100 reserved cases across four intent classes and all 31 category identifiers in each split. Development v2 records 48/48 for both parsers; this is development fit, not test evidence or an LLM advantage.

What blocks completion. The protocol requires a genuine independent student annotation pass, an exposure declaration, agreement measurement and blinded adjudication before reserved evaluation. AI cannot supply that pass in the student's name. Preserve the parser commitment dated 29 September 2026 at 05:32 UTC. The 3 October runtime check found no Ollama executable and a refused localhost connection, so new Llama inference is stopped. Verify the recorded Llama 3.1 8B runtime, complete the prerequisites and full freeze, then evaluate once and report parser accuracy before ranking results. If H1 or H2 fails, report that natural-language feedback did not demonstrate an expressiveness advantage under this protocol and identify why.

For the supervisor. Confirm individual scope approval, the COMP 9500 path, the original NI margin reading, AI disclosure and the final deadline. The official BCIT presentation template is now supplied. The plan includes a 23 October readiness decision and a complete paper draft by 27 November. The proposed remaining schedule is 5 October to 18 December: 135 future student hours over 11 weeks, with Wednesday 10:00 am Vancouver meetings, Tuesday 10:00 am updates and minutes afterwards. It is a plan, not a completed activity log; the exact December deadline remains unconfirmed.

Verification and limits. September 29 audit results are retained as historical records. October checks must be read from their separately dated evidence and reports. The later numerical audit reproduced all 25,950 top-10 lists, with 2,009 full-ranking hash differences and 5,190 request-score hash differences still unresolved. Reserved Experiment 2 replay remains blocked. Windows 11 and Colab execution are unverified. Preparing this package does not establish student proficiency, completed hours, supervisor approval, publication readiness or final course completion.